

Distributed and Cloud Computing

CYEN405 / CSC452 / CSC552

Parallel and Distributed Computing with MATLAB

Ben Drozdenko, Ph.D.
Assistant Professor, Cyber Engineering & CS
September 11, 2018

Lecture Outline

- Parallel Computing Toolbox™ (PCT) and MATLAB® Distributed Computing Server™ (MDCS)
 - Starting a Parallel Pool using `parpool`
- Task-Parallel Jobs with PCT
 - Embarrassingly Parallel Tasks using `parfor`
 - Setting up a Simple Independent Job using `createJob`
- Data-Parallel Jobs with PCT
 - Single Program, Multiple Data using `spmd`
 - Distributed Datasets and Operations
 - Passing Messages
 - Setting up a Communicating Job

MathWorks® Products

- MATLAB
 - High-level language and interactive environment for numerical computation, visualization, and programming. Language, tools, and built-in math functions to explore multiple approaches and reach a solution faster.
 - Used for a range of applications, including signal processing and communications, image and video processing, control systems, test and measurement, computational finance, and computational biology.
- Parallel Computing Toolbox™ (PCT)
 - Solve computationally-intensive and data-intensive problems using multicore processors, GPUs, and computer clusters.
 - High-level constructs—parallel for-loops, special array types, and parallelized numerical algorithms—to parallelize MATLAB® applications without CUDA or MPI programming.
- MATLAB® Distributed Computing Server™ (MDCS)
 - Run computationally intensive MATLAB® programs and Simulink® models on computer clusters, clouds, & grids. Develop your program or model on a multicore desktop computer using PCT and then scale up to many computers by running it on MDCS.
 - Server includes a built-in cluster job scheduler and provides support for commonly used third-party schedulers (e.g. Platform LSF, MS Windows HPCS, PBS TORQUE, etc.).
- Source: MathWorks® Product Page: <http://www.mathworks.com/products/>

CYEN405/CSC452/CSC552 Fall 2018

3

Starting Parallel Pool with **parpool**

- To get an estimate of the #computational cores on local machine:
`>> maxNumCompThreads`
- To open a pool of 4 workers, type at the command prompt:
`>> parpool(4)`
- This command starts 4 new instances of MATLAB, which are available for computations as part of a resource pool.
- In Windows, type Ctrl-Alt-Del to start Task Manager. Select Processes. In total, there are now 4 instances of MATLAB.exe running on the system—4 as part of the pool + original instance.
- To query the current pool:
`>> p = gcp; p.NumWorkers`
- A parallel pool automatically starts when you execute a parallel language construct that runs on a pool, such as **parfor** or **spmd**.
- When finished, to close pool, type:
`>> delete(gcp)`

CYEN405/CSC452/CSC552 Fall 2018

4

Lecture Outline (2)

- Parallel Computing Toolbox™ (PCT) and MATLAB® Distributed Computing Server™ (MDCS)
 - Starting a Parallel Pool using `parpool`
- Task-Parallel Jobs with PCT
 - Embarrassingly Parallel Tasks using `parfor`
 - Setting up a Simple Independent Job using `createJob`
- Data-Parallel Jobs with PCT
 - Single Program, Multiple Data using `spmd`
 - Distributed Datasets and Operations
 - Passing Messages
 - Setting up a Communicating Job

CYEN405/CSC452/CSC552 Fall 2018

5

Task-Parallel Jobs with PCT

- Parallel For Loops with `parfor` keyword
 - Each `worker` runs independently
 - Ideally suited for embarrassingly-parallel tasks
- Parallel Jobs using `createJob` function
 - Assign specific tasks to run on workers
- Specific rules for usage
 - No communication between loop iterations

CYEN405/CSC452/CSC552 Fall 2018

6

3

Example #1: Birthday Paradox

- What is the probability that in a group of 30 randomly selected individuals, at least two of the individuals will share the same birthday?

Assuming independent events, $p_{bday} = 1 - \frac{30! \cdot \binom{365}{30}}{365^{30}} \approx 70.6\%$

- MATLAB® code for one trial:

```
function match = birthday(groupSize)
bdays = randi(365, groupSize, 1);
bdays = sort(bdays);
match = any(diff(bdays) == 0);
```

- Code to run many trials sequentially (“brute-force algorithm”):

```
function prob = runBirthday(numtrials, groupsize)
matches = false(1,numtrials);
for trial = 1:numtrials
    matches(trial) = birthday(groupsize);
end
prob = sum(matches)/numtrials;
```

CYEN405/CSC452/CSC552 Fall 2018

7

Parallel For Loops with **parfor**

- Use the keyword **parfor** to make any **for** loop into a parallel loop that runs the *independent* iterations on different workers
- Code to run many trials in parallel:

```
function prob = pRunBirthday(numtrials, groupsize)
matches = false(1,numtrials);
parfor trial = 1:numtrials
    matches(trial) = birthday(groupsize);
end
prob = sum(matches)/numtrials;
```

CYEN405/CSC452/CSC552 Fall 2018

8

Time Code using **tic** and **toc**

- With pool still open, time parallel version

```
>> tic; p = pRunBirthday(1e5,30), toc
```

- Close pool, and time sequential version

```
>> delete(gcp);
```

```
>> tic; p = runBirthday(1e5,30), toc
```

- On my local machine

N_{trials}	1e5	1e6
Sequential	0.40 sec	4.00 sec
Parallel, $N_{workers}=4$	0.24 sec	1.81 sec
Speedup	1.67X	2.2X

Setup a Parallel Job using **createJob**

- Observe the behavior of a parallel for loop from the following output:

```
>> parfor i=1:10, disp(i); end
```

- Parfor allows for little control over parallel execution. To assign specific tasks to each worker, create an independent job instead:

- Connect to a cluster:

```
>> cluster = parcluster('local');
```

- Create an independent job:

```
>> job = createJob(cluster);
```

Setup a Parallel Job using `createJob` (cont.)

3. Create many tasks for the job to handle (you could use a `for` loop or a `while` loop for this).


```
>> createTask(job, @runBirthday, 1, {1e5,5});
>> createTask(job, @runBirthday, 1, {1e5,10});
>> createTask(job, @runBirthday, 1, {1e5,15});
>> createTask(job, @runBirthday, 1, {1e5,20});
>> createTask(job, @runBirthday, 1, {1e5,25});
>> createTask(job, @runBirthday, 1, {1e5,30});
```
4. Submit the job. Optionally, wait for it to finish.


```
>> submit(job); wait(job, 'finished');
```
5. Retrieve the results.


```
>> results = fetchOutputs(job);
>> results{end,1}
>> r = cell2mat(results); mean(r)
```
6. When finished, delete job & clear object.


```
>> delete(job); clear job;
```

CYEN405/CSC452/CSC552 Fall 2018

11

Example #2: Gene Matching

```
function results = pargenematchsol()
searchSeq = repmat('gattaca', 1, 10);
numTasks = 2;
numBases = 7048095;
cluster = parcluster('local');
job = createJob(cluster);
[startValues, endValues] = splitDataset(numBases, numTasks);
offsetLeft = floor(length(searchSeq)/2);
if mod(length(searchSeq),2) == 0
    offsetRight = offsetLeft - 1;
else
    offsetRight = offsetLeft;
end
startValues(2:end) = startValues(2:end) - offsetLeft;
endValues(1:end-1) = endValues(1:end-1) + offsetRight;
for tasknum = 1:numTasks
    createTask(job, @genematch, 2, {searchSeq, 'gene.txt', ...
        startValues(tasknum), endValues(tasknum)} );
end
```

CYEN405/CSC452/CSC552 Fall 2018

12

Example #2: Gene Matching (cont.)

```

submit(job); % Submit and Wait for Results
wait(job, 'finished');
results = fetchOutputs(job); % Report the results
results = cell2mat(results); % Return absolute position
[~,idx] = max(results(:,1));
bpm = results(idx,1);
msi = results(idx,2)+startValues(idx)-1;

function [startValues, endValues] = splitDataset(numTotalElements, numTasks)
numPerTask = repmat(floor(numTotalElements/numTasks), 1, numTasks);
leftover = rem(numTotalElements, numTasks);
numPerTask(1:leftover) = numPerTask(1:leftover) + 1;
endValues = cumsum(numPerTask);
startValues = [1 endValues(1:end-1) + 1];

```

CYEN405/CSC452/CSC552 Fall 2018

13

Example #2: Gene Matching (cont.)

```

function [bestPctMatch,matchStartIdx]=genematch(searchSeq,file,startIdx,endIdx)
fid = fopen(file, 'rt');
geneSeq = fscanf(fid, '%c');
fclose(fid);
if nargin < 3, startIndex = 1; end
if nargin < 4, endIndex = length(geneSeq); end
[bestPctMatch,matchStartIdx]=findsubstr(geneSeq(startIdx:endIdx),searchSeq);

function [bestPctMatch,matchStartIdx]=findsubstr(baseString,searchString)
bestPctMatch = 0;
matchStartIdx = 0;
for startIdx = 1:(length(baseString)-length(searchString)+1)
    currentSection = baseString(startIdx:startIdx+length(searchString)-1);
    pctMatch = nnz(currentSection==searchString)/length(searchString);
    if pctMatch >= bestPctMatch
        bestPctMatch = pctMatch;
        matchStartIdx = startIdx;
    end
end

```

CYEN405/CSC452/CSC552 Fall 2018

14

Lecture Outline (3)

- Parallel Computing Toolbox™ (PCT) and MATLAB® Distributed Computing Server™ (MDCS)
 - Starting a Parallel Pool using `parpool`
- Task-Parallel Jobs with PCT
 - Embarrassingly Parallel Tasks using `parfor`
 - Setting up a Simple Independent Job using `createJob`
- Data-Parallel Jobs with PCT
 - Single Program, Multiple Data using `spmd`
 - Distributed Datasets and Operations
 - Passing Messages
 - Setting up a Communicating Job

CYEN405/CSC452/CSC552 Fall 2018

15

Data-Parallel Jobs with PCT

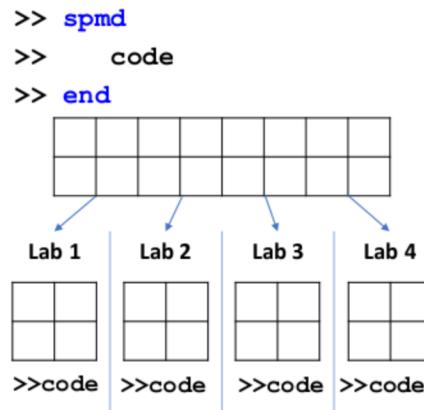
- Single Program, Multiple Data with `spmd`
 - Self-identification using `labindex` and `numlabs`
 - Types of arrays—replicated, variant, private
 - Composite data type
- Distributed Datasets and Operations
- Passing Messages
 - Practical Considerations
- Setting up Communicating Jobs using `createCommunicatingJob`
 - Assign one task to run on all labs
 - A lab can pass messages to other labs

CYEN405/CSC452/CSC552 Fall 2018

16

Single Program, Multiple Data: **spmd**

```
>> spmd
>> labindex
>> numlabs
>> end
Lab 1:
    1
    4
Lab 2:
    2
    4
Lab 3:
    3
    4
Lab 4:
    4
    4
```



CYEN405/CSC452/CSC552 Fall 2018

17

Types of Arrays on Labs in SPMD

- Replicated Array

```
>> spmd
>> x = 5;
>> end
```

Lab 1	x = 5
Lab 2	x = 5
Lab 3	x = 5
Lab 4	x = 5

Lab 1	y = 0.3246
Lab 2	y = 0.2646
Lab 3	y = 0.8847
Lab 4	y = 0.8939

Lab 1	
Lab 2	z = 7
Lab 3	
Lab 4	

CYEN405/CSC452/CSC552 Fall 2018

18

Composite Class & Reductions

- From the MATLAB client, all these three types of arrays show up in the workspace as a Composite data type.


```
>> class(y)
      ans = composite
```

 - Use the curly braces to extract the contents at any lab index.


```
>> y{3}
```
- Use a global operation to combine the results from all the labs. This performs a Reduction & Broadcast, which turns a Composite variable into a replicated array.


```
>> spmd
      >> a = gcat(y); % Global concatenation
      >> s = gplus(y); % Global summation
      >> m = gop(@max,y); % Global maximum
      >> end
```

 - Or, specify a lab index to perform a Reduction and turn it into a private array:


```
>> spmd
      >> a1 = gcat(y,1,1);
      >> s1 = gplus(y,1);
      >> m1 = gop(@max,y,1);
      >> end
```

CYEN405/CSC452/CSC552 Fall 2018

19

Distributed Datasets and Operations

- Use Distributed Data Type from Client.


```
vars = load('airportdata');
dlat = distributed(vars.lat);
dlon = distributed(vars.long);
```
- OR: Use Codistributed Data Type from Labs.


```
spmd
    vars = load('airportdata');
    clat = codistributed(vars.lat);
    clong = codistributed(vars.long);
end
```
- Use `gather` function to convert to replicated array.


```
>> allat = gather(dlat);
      % Or, specify a lab index to convert to a private array.
      >> allat1 = gather(dlat, 1);
```
- Also, use `getLocalPart` to convert to a variant array.


```
>> loclat = getLocalPart(dlat);
```

CYEN405/CSC452/CSC552 Fall 2018

20

10

Example #3: Airport Distances

```

spmd
    R = 3963.2; % in miles
    vars = load('airportdata');
    lat = codistributed(vars.lat);
    long = codistributed(vars.long);
    lat = 90 - lat;
    long = 360 + long;
    x = R * sind(lat) .* cosd(long);
    y = R * sind(lat) .* sind(long);
    z = R * cosd(lat);
    coords = [x y z]';
    dotprod = coords' * coords;
    mag = sqrt(sum(coords.^2));
    angles = min(dotprod ./ (mag' * mag), 1);
    dist = R * acos(angles); % Arc length
end

```

CYEN405/CSC452/CSC552 Fall 2018

21

Passing Messages

- To send a variable **x** to another lab:

```
>> labSend(x, dest_labindex);
```

- To receive a variable **x** from another lab:

```
>> x = labReceive(src_labindex);
```

- To see whether a lab is ready to receive data:

```
>> isReady = labProbe(dest_labindex);
```

- To broadcast a variable **x** to all other labs:

```
>> x = labBroadcast(src_lab,x);
```

- To synchronize all the labs:

```
>> labBarrier;
```

CYEN405/CSC452/CSC552 Fall 2018

22

Passing Messages: Practical Considerations

```

spmd
    switch labindex
        case 1
            x1 = labindex * ones(1, 5); % Create local data
            x2 = labReceive(2); % Receive data from peer
            labSend(x1, 2); % Send data to peer
            y = x2; % Return peer's data
        case 2
            x2 = labindex * ones(1, 5); % Create local data
            x1 = labReceive(1); % Receive data from peer
            labSend(x2, 1); % Send data to peer
            y = x1; % Return peer's data
    end

```

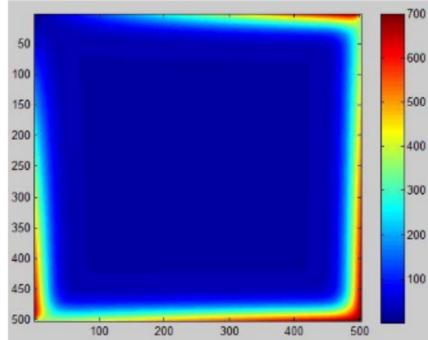
Deadlock! Use `labSendReceive` function instead to exchange data between labs.

```

spmd
    switch labindex
        case 1
            x1 = labindex * ones(1, 5); % Create local data
            x2 = labSendReceive(2, 2, x1); % Exchange data with lab 2
            y = x2; % Return peer's data
        case 2
            x2 = labindex * ones(1, 5); % Create local data
            x1 = labSendReceive(1, 1, x2); % Exchange data with lab 1
            y = x1; % Return peer's data
    end

```

Example #4: Parallel Heat Equation



- Get a matrix representing the temperature at each point of a 2D square plate of length L & diffusivity C
- For example, solve for the temperature on a 3m-by-3m copper plate after 40 seconds have elapsed, using 500 time steps of 80 ms each. Thermal diffusivity of copper is $1.13e-4 \text{ m}^2/\text{s}$.

Example #4: Sequential Heat Equation

```

function U = heateq(k, n, Ts, L, c)
ms = L / n;
if Ts > (ms^2/2/c), error('Selected time step is too large.');?>
U = initialTempDistrib(n);
north = 1:n;
south = 3:(n + 2);
curr = 2:(n + 1);
east = 3:(n + 2);
west = 1:n;
for iter = 1:k
    U(curr, curr) = U(curr, curr) + c * Ts/(ms^2) * (U(north, curr) + ...
        U(south, curr) - 4*U(current, curr) + U(curr, east) + U(curr, west));
end

function U = initialTempDistrib(n)
U = 23*ones(n + 2);
U(1, :) = (1:(n + 2))*700/(n + 2);
U(end, :) = ((1:(n + 2)) + (n + 2))*700/2/(n + 2);
U(:, 1) = (1:(n + 2))*700/(n + 2);
U(:, end) = ((1:(n + 2)) + (n + 2))*700/2/(n + 2);

```

CYEN405/CSC452/CSC552 Fall 2018

25

Setting up Communicating Jobs

1. Connect to a cluster.
`>> cluster = parcluster('local');`
2. Create a communicating job.
`>> job = createCommunicatingJob(cluster, 'Type', 'SPMD');`
3. Create a repeating task for the job to handle.
`>> createTask(job, @parheateqn, 1, ...
 {1e3, 500, 0.08, 3, 1.13e-4});`
4. Set a range for the number of workers needed.
`>> set(job, 'NumWorkersRange', [3 3]);`
5. Submit the job. Optionally, wait for it to finish.
`>> submit(job); wait(job, 'finished');`
6. Retrieve the results.
`>> results = fetchOutputs(job);
>> U = cell2mat(results');
>> imagesc(U)`
7. When finished, clean up. Delete job & clear object.
`>> delete(job); clear job;`

CYEN405/CSC452/CSC552 Fall 2018

26

Example #4: Parallel Heat Equation

```

function U = parheateqn(k, n, Ts, L, c)
ms = L / n;
if (Ts>(ms^2/2/c)), error('Selected time step is too large.');?>
Uinit = initialTempDistrib(n);
parts = codistributor1d.defaultPartition(n+2);
numLocalCols = parts(labindex);
leftColInd = sum(parts(1:labindex - 1)) + 1;
rightColInd = leftColInd + numLocalCols - 1;
U = Uinit(:, leftColInd:rightColInd);
if (labindex > 1), U = [zeros(n+2, 1) U]; end
if (labindex < numlabs), U = [U zeros(n+2, 1)]; end
if (labindex == 1) || (labindex == numlabs)
    numLocalCols = numLocalCols - 1;
end
rightNeighbor = mod(labindex, numlabs) + 1;
leftNeighbor = mod(labindex - 2, numlabs) + 1;
north = 1:n;
south = 3:n + 2;
currRow = 2:n + 1;
currCol = 2:numLocalCols + 1;
east = 3:numLocalCols + 2;
west = 1:numLocalCols;

```

CYEN405/CSC452/CSC552 Fall 2018

27

Example #4: Parallel Heat Equation (cont.)

```

for iter = 1:k
    rightBoundary = labSendReceive(leftNeighbor,rightNeighbor,U(:,2));
    leftBoundary = labSendReceive(rightNeighbor,leftNeighbor,U(:,end-1));
    if (labindex > 1), U(:, 1) = leftBoundary; end
    if (labindex < numlabs), U(:, end) = rightBoundary; end
    % Update grid for current iteration
    U(currRow,currCol) = U(currRow,currCol) + ...
        c*Ts/(ms^2)*(U(north,currCol) + U(south,currCol) - ...
        4*U(currRow,currCol) + U(currRow,east) + U(currRow,west)));
end
% Combine parts from all labs into a single matrix stored on lab 1
U = gcat(U(currRow, currCol), 2, 1);

function U = initialTempDistrib(n)
U = 23*ones(n + 2);
U(1, :) = (1:(n + 2))*700/(n + 2);
U(end, :) = ((1:(n + 2)) + (n + 2))*700/2/(n + 2);
U(:, 1) = (1:(n + 2))*700/(n + 2);
U(:, end) = ((1:(n + 2)) + (n + 2))*700/2/(n + 2);

```

CYEN405/CSC452/CSC552 Fall 2018

28

Summary: Problem Types

	Interactive	Batch
Task-Parallel	parpool parfor	createJob createTask
Data-Parallel	parpool spmd	createCommunicatingJob createTask